

TIP 0002: Contiguity Argument for Memory Consistency

TIP	0002
authors:	Alan Szeplieniec
title:	Multi-Table Subset Argument for Memory Consistency
status:	draft
created:	2022-08-25
issue tracker:	
pdf:	tip-0002.pdf

Abstract. In the current specification, the memory-like tables `RamTable`, `JumpStackTable`, and `OpStackTable` do not satisfy memory-consistency. Specifically, they are vulnerable to Yuncong’s attack, which exploits the unverified and thus possibly-incorrect sorting in these tables. This TIP addresses one part of the issue by introducing a new table argument, the *multi-table subset argument*. It establishes that certain values in different tables have a matching representative in a given lookup table. Applied to the present context, this technique establishes that every clock jump is positive. This note is a companion to TIP-0001, which introduces an new argument to establish the contiguity of regions of constant memory pointer. Together, TIP-0001 and TIP-0002 fix the memory consistency issue.

Introduction

How to establish that a clock jump is directed forward (as opposed to direct backward, which would indicate malicious behavior)? One strategy is to show that the *difference*, *i.e.*, the next clock cycle minus the current clock cycle, is itself a clock cycle. Recall that the Processor Table’s clock cycles run from 0 to $T - 1$. Therefore, valid differences belong to $\{2, \dots, T - 1\}$ and invalid ones to $\{p - T + 1, \dots, p - 2\}$.

Standard subset arguments can show that the clock jump differences are elements of the Processor Table’s clock cycle column. However, it is cumbersome to repeat this argument for three separate tables.

I present here a Polynomial IOP / RAP hybrid argument for showing that all clock jump differences in all three memory-like tables live also in the Processor Table’s `clk` column. It introduces one extension column in each memory-like table; one extension column in Processor Table; and three committed polynomials that do not correspond to any columns.

Intuition

Let $f(X)$ denote the polynomial

$$f(X) = \left(\prod_{\delta_r} (X - \delta_r) \right) \cdot \left(\prod_{\delta_o} (X - \delta_o) \right) \cdot \left(\prod_{\delta_j} (X - \delta_j) \right)$$

where the δ_r , δ_o , and δ_j denote the clock jump differences in the Ram Table, OpStack Table, and Jump Stack Table, respectively.

This polynomial can be (and is) evaluated in a scalar α supplied by the verifier through three running products. (All three running products use the same challenge α .) However, it is additionally committed to. The point α is used to verify that the committed polynomial $[f(X)]$ does indeed correspond to the product of running product columns.

Expand the Processor Table with a new column `cji` (“clock jump indicator”) which takes the value 1 if the `clk` value is also the value of some clock jump, and 0 otherwise. The extension column computes the value of the polynomial

$$h(X) = \prod_{\delta} (X - \delta)$$

in some scalar β supplied by the verifier, where δ ranges over the values of `clk` in indicated rows.

Note that the roots of $f(X)$ correspond to roots of $h(X)$. However, the roots occur with various multiplicities in $f(X)$ and with multiplicity 1 in $h(X)$. To harmonize the multiplicities, we need the formal derivative $f'(X)$ of $f(X)$ to cancel all the square factors. Specifically, let $g(X) = \gcd(f(X), f'(X))$ then $h(X)g(X) = f(X)$ and to establish the correct gcd relation we need Bézout coefficients $a(X)$ and $b(X)$ for the Bézout relation

$$a(X)f(X) + b(X)f'(X) = g(X).$$

Detailed Description

Memory-like Tables

Here are the constraints for the RAM Table. The constraints for other two tables are analogous and are therefore omitted. The constraints are relative to the indeterminate X which anticipates the verifier’s challenge α via the implicit substitution $X \mapsto \alpha$.

Use `mp` to abstractly refer to the memory pointer. The extension column starts with a random initial `rpI` and no boundary constraints. Before the substitution $X \mapsto \alpha$, `rpI` is a uniformly random polynomial of degree at most 3 in X .

The transition constraint enforces the accumulation of a factor $(X - \text{clk}^* + \text{clk})$ whenever there is a clock jump and the memory pointer is the same. If there is

no clock jump or the memory pointer is unchanged, the same running product is carried to the next row. The transition constraint is

$$(\text{mp}^* - \text{mp}) \cdot (\text{clk}^* - \text{clk} - 1) \cdot (\text{rp}^* - \text{rp} \cdot (X - \text{clk}^* + \text{clk})) + (\text{mp}^* - \text{mp}) \cdot (\text{clk}^* - \text{clk}) \cdot (\text{rp}^* - \text{rp}) + (1 - (\text{mp}^* - \text{mp}) \cdot \text{di}) \cdot (\text{rp}^* - \text{rp})$$

Note that di is the difference inverse of the Ram Table but for the other two tables this factor can be dropped.

The running product has a terminal rp_T but a trivial terminal constraint: $\text{rp} - \text{rp}_T$.

The prover computes $\text{rp}_T(X)$ symbolically for every memory-like table. The polynomial $f(X)$ is the product of all such symbolic terminals.

Processor Table

The Processor Table also has a running product extension column rp . In the first row it is unconstrained and set to a random initial $\text{rp}_I \in \mathbb{F}$. In every row it is either updated with a factor $(X - \text{clk})$ or not – the prover knows what to do. The verifier only needs to verify the transition constraint $(\text{rp}^* - \text{rp}) \cdot (\text{rp}^* - \text{rp} \cdot (X - \text{clk}))$. The terminal constraint is likewise $(\text{rp}_T - \text{rp}) \cdot (\text{rp}_T - \text{rp} \cdot (X - \text{clk}))$.

The formal indeterminate X here anticipates the verifier's second challenge β . The prover computes rp_T symbolically before β is supplied – in fact, this symbolic terminal is the polynomial $h(X)$.

Square-Freeness Argument

All memory-like tables' clock jump differences are members of the Processor Table's clk column if $h(X)$ is the maximal square-free divisor of $f(X)$. This relation is implied by three points: - $f'(X)$ is the formal derivative of $f(X)$ - $g(X)$ is the greatest common divisor of $f(X)$ and $f'(X)$ - $h(X) = f(X)/g(X)$.

To see this, let $f(X) = \prod_i f_i(X)^{m_i}$. Then $f'(X) = \sum_i m_i f_i(X)^{m_i-1} f'_i(X) \prod_{j \neq i} f_j(X)$. And $g(X) = \prod_i f_i(X)^{m_i-1}$ because every factor in the product appears in all terms of $f'(X)$, but there will be one term that limits the multiplicity to $m_i - 1$. And from this it follows that $h(X) = f(X)/g(X)$.

Note that it is possible to eliminate $g(X) = f(X)/h(X)$ from the above bullet points. Therefore, two relations remain to be shown: the formal derivative, and the gcd.

Greatest Common Divisor

Let $a(X)$ and $b(X)$ be randomized Bézout coefficients such that

$$a(X)f(X) + b(X)f'(X) = g(X) \ .$$

Specifically, the prover computes coefficients $a^*(X)$ and $b^*(X)$ using the extended Euclidean algorithm. Then he samples a random factor $k(X)$ of degree at most 2 and sets $a(X) = a^*(X) + k(X)f'(X)$ and $b(X) = b^*(X) - k(X)f(X)$.

The verifier probes this identity of polynomials in a randomly chosen point $z_1 \xleftarrow{\$} \mathbb{F}$.

Formal Derivative

The formal derivative of a polynomial $f(X) = \sum_{i=0}^d c_i X^i$ is $f'(X) = \sum_{i=1}^d c_i i X^{i-1}$. The correct relation between two polynomial oracles $[f(X)]$ and $[f'(X)]$ can be verified by probing the relation $X \cdot f'(X) = (N_d \circ f)(X)$, where $N_d(X) = \sum_{i=1}^d i X^i$ is the natural arithmetic polynomial up to some degree bound $d \geq \deg(f(X))$, and where \circ denotes the Hadamard (coefficient-wise) product.

Note that the natural arithmetic polynomial can be evaluated efficiently via the following recursive relation.

$$\begin{aligned}
N_d(X) &= \sum_{i=1}^d i X^i = \sum_{i=0}^{\frac{d+1}{2}-1} i X^i + \sum_{i=0}^{\frac{d+1}{2}-1} \left(\frac{d+1}{2} + i\right) X^{\frac{d+1}{2}+i} \\
&= \sum_{i=0}^{\frac{d+1}{2}-1} i X^i + X^{\frac{d+1}{2}} \cdot \left(\frac{d+1}{2} \sum_{i=0}^{\frac{d+1}{2}-1} X^i + \sum_{i=0}^{\frac{d+1}{2}-1} i X^i \right) \\
&= \left(1 + X^{\frac{d+1}{2}}\right) \sum_{i=0}^{\frac{d+1}{2}-1} i X^i + \frac{d+1}{2} \cdot X^{\frac{d+1}{2}} \cdot \frac{X^{\frac{d+1}{2}} - 1}{X - 1} \\
&= \left(1 + X^{\frac{d+1}{2}}\right) N_{\frac{d+1}{2}-1}(X) + \frac{d+1}{2} \cdot \frac{X^{d+1} - X^{\frac{d+1}{2}}}{X - 1}
\end{aligned}$$

Therefore, no further polynomial oracles are required to verify the formal derivative relation beyond those that are required to verify the Hadamard relation. For verifying the Hadamard relation, the §3.5 of the Claymore paper provides a solution, which is summarized below in unrolled form.

- The verifier samples $\alpha \xleftarrow{\$} \mathbb{F} \setminus \{0\}$ and sends η to the prover.
- The prover computes the coefficient vector \mathbf{c} of $N_d(\alpha \cdot X) \cdot f(X^{-1}) \cdot X^d$ and the polynomial $\bar{h}(X) = \sum_{i=0}^{2d} \bar{h}_i X^i$, where the i th coefficient is $\bar{h}_i = \frac{c_i}{\gamma^{d-\gamma^i}}$ and $\bar{h}_d \xleftarrow{\$} \mathbb{F}$. The prover sends $\bar{h}(X)$ to the verifier along with the implicit claim that its degree is bounded by $2d$. Here $\gamma \in \mathbb{F}$ is some system-wide parameter with a sufficiently large multiplicative order.

- The verifier probes the polynomial identity $\bar{h}(X) \cdot \gamma^d - \bar{h}(\gamma X) + \alpha \cdot f'(\alpha) \cdot X^d = N_d(\alpha \cdot X) \cdot f(X^{-1}) \cdot X^d$ in a random point $z \xleftarrow{\$} \mathbb{F} \setminus \{0\}$.

Putting Everything Together

1. The prover computes $f(X)$, $h(X)$, $f'(X)$, $a(X)$, $b(X)$ and all base columns.
2. The prover commits to $f(X)$, $a(X)$, $b(X)$ and to all base columns.
3. The verifier supplies uniformly random challenges α, β, z .
4. The prover computes the extension columns, and $\bar{h}(X)$, and commits to them. He also reveals the terminals.
5. Let α be the challenge for the Processor Table's running product such that $h(\alpha) = \mathbf{rp}_T$. The prover reveals $a(\alpha), b(\alpha), f(\alpha)$. Note that $f'(\alpha) = \frac{a(\alpha)f(\alpha) - f(\alpha)/h(\alpha)}{b(\alpha)}$.
6. The prover reveals $f(z^{-1})$, $\bar{h}(z)$ and $\bar{h}(\gamma z)$.
7. The verifier probes the polynomial identity $\bar{h}(X) \cdot \gamma^d - \bar{h}(\gamma X) + \alpha \cdot f'(\alpha) \cdot X^d = N_d(\alpha \cdot X) \cdot f(X^{-1}) \cdot X^d$ in $X = z$ with the revealed values.
8. Let β be the challenge for computing the running products of the memory-like tables, resulting in three terminals which, multiplied together, give $f(\beta)$. Let $i_f(X)$ be the lowest-degree interpolant through $(\alpha, f(\alpha)), (\beta, f(\beta)), (z^{-1}, f(z^{-1}))$, and $i_{\bar{h}}(X)$ the lowest-degree interpolant through $(z, \bar{h}(z)), (\gamma z, \bar{h}(\gamma z))$. Add to the nonlinear combination:
 - the polynomials $f(X), a(X), b(X), \bar{h}(X)$ of degrees at most $T + 9$, $T + 8$, $T + 9$, and $2T + 18$, respectively.
 - the quotient $\frac{f(X) - i_f(X)}{(X - \alpha)(X - \beta)(X - z)}$ of degree at most $T + 6$
 - the quotient $\frac{a(X) - a(\alpha)}{X - \alpha}$ of degree at most $T + 7$
 - the quotient $\frac{b(X) - b(\alpha)}{X - \alpha}$ of degree at most $T + 8$.
 - the quotient $\frac{\bar{h}(X) - i_{\bar{h}}(X)}{(X - z)(X - \gamma z)}$ of degree at most $2T + 16$.

The reason why $f(X)$ has degree at most $T + 9$, not $3T + 9$ as one might expect, is because any instruction can affect only one memory-like table. As a result, every cycle generates at most one clock jump.

Security

The claim is that within each contiguous region of constant memory pointer, the clock cycle column is sorted in ascending order. The AIR verifies increments by one. Therefore the claim boils down to this: increments by more than one within the same contiguous region – *clock jumps* – increase the clock cycle column by some value in $\{2, \dots, T\}$.

The entire technique reduces this claim to an identity of polynomials. Specifically, the polynomial $\prod_{\delta} (X - \delta)$, where δ ranges over all clock jump differences, after removing all factors with multiplicity greater than one, must be equal to the polynomial $\prod_i (X - \text{clk}_i)$, where the product ranges over all clock cycles that correspond to some clock jump difference.

Completeness

Completeness follows from construction, except in special cases where α . Note that if α coincides with one of the roots of the initials of one of the memory-like tables, then $f(X) = f'(X) = 0$ but this degeneration does not affect completeness. However, if two of the initials have a nontrivial common divisor, then $f'(X)$ will contain a factor that $h(X)$ does not contain. In this case the prover fails.

What is the probability that the initials have non-trivial common factors? The initials have degree at most one and we may assume without loss of generality that they have degree exactly one because otherwise they cannot have non-trivial factors to begin with. Polynomials of degree 1 that share non-trivial factors must be multiples of each other. There is 1 way to sample the leading coefficient of the second initial such that it is a multiple of the first initial, and there are 2 ways to sample the leading coefficient of the third initial such that it is a multiple of either the first or the second. By the union bound, the probability of the initials sharing a non-trivial factor is therefore $3/|\mathbb{F}|$. This quantity is the completeness error. \square

Soundness

The square-freeness argument generates a false positive when $a(X)f(X) + b(X)f'(X) \neq g(X)$ except in the point $X = \alpha$ where it was sampled. The degree of this polynomial identity is $2T + 17$, and so by the Schwarz-Zippel lemma the probability of false positive is at most $(2T + 17)/|\mathbb{F}|$.

The formal derivative argument boils down to a Hadamard product, whose soundness error is bounded by $3(T + 9)/|\mathbb{F}|$. By the union bound, the soundness error of the entire construction is bounded by $(5T + 26)/|\mathbb{F}|$. \square

Zero-Knowledge

The prover releases 4 terminals $\mathbf{rp}_T^{(r)}$, $\mathbf{rp}_T^{(o)}$, $\mathbf{rp}_T^{(j)}$, $\mathbf{rp}_T^{(p)}$, which are the running products of the three memory-like tables and the Processor Table, respectively. In addition to that, the prover also releases the evaluations $f(\alpha)$, $a(\alpha)$, $b(\alpha)$, $f(z^{-1})$, $\bar{h}(z)$, $\bar{h}(\gamma z)$.

For every tuple $\mathbf{rp}_T^{(r)}$, $\mathbf{rp}_T^{(o)}$, $\mathbf{rp}_T^{(j)}$, $f(\alpha)$, $f(z^{-1})$, there are consistent random initials $\mathbf{rp}_I^{(r)}$, $\mathbf{rp}_I^{(o)}$, $\mathbf{rp}_I^{(j)}$.

- The mapping $\mathbf{rp}_{I,0}^{(o)} \mapsto \mathbf{rp}_T^{(o)}$ is affine and the coefficient is the product $\prod(\beta - \delta^{(o)})$ where $\delta^{(o)}$ ranges over all clock jumps in the OpStack table. This coefficient is zero with negligible probability.
- The mapping $\mathbf{rp}_{I,0}^{(j)} \mapsto \mathbf{rp}_T^{(j)}$ is affine and the coefficient is the product $\prod(\beta - \delta^{(j)})$ where $\delta^{(j)}$ ranges over all clock jumps in the JumpStack table. This coefficient is zero with negligible probability.

- The mapping $(\mathbf{rp}_{I,0}^{(r)}, \mathbf{rp}_{I,1}^{(r)}, \mathbf{rp}_{I,2}^{(r)}, \mathbf{rp}_{I,3}^{(r)}) \mapsto (f(\beta), f(\alpha), f(z^{-1}), f(\gamma^{-1}z^{-1}))$ is invertible with high probability because it is affine and its coefficient matrix is Vandermonde. A singular Vandermonde matrix requires $\alpha, \beta, z^{-1}, z, \gamma z$ to coincide to have low order, which occurs with negligible probability.
 - The mapping $f(\beta) \mapsto \mathbf{rp}_T^{(r)}$ is invertible with high probability because the factor $\mathbf{rp}_T^{(o)} \mathbf{rp}_T^{(j)}$ is fixed and nonzero with overwhelming probability.
 - The mapping $f(\gamma^{-1}z^{-1}) \mapsto \bar{h}(\gamma z) = N_d(\alpha \cdot \gamma z) \cdot f(\gamma^{-1}z^{-1}) \cdot \gamma^d z^d$ is linear and the coefficient is nonzero with overwhelming probability.

For every set of initials fixed so far, for every evaluation $\bar{h}(z)$ there is a consistent random coefficient \bar{h}_d . The mapping $\bar{h}_d \mapsto \bar{h}(z)$ is affine and the coefficient is z^d , which is invertible with overwhelming probability.

For every terminal $\mathbf{rp}_T^{(p)}$ there is a consistent initial $\mathbf{rp}_I^{(p)}$ because the mapping $\mathbf{rp}_I^{(p)} \mapsto \mathbf{rp}_T^{(p)}$ is linear and the coefficient is given by the product $\prod(\alpha - \text{clk}_i)$ which ranges over all relevant clock cycles. This coefficient is nonzero with overwhelming probability.

For every set of initials, and for every evaluation $a(\alpha)$ there is a consistent Bézout randomizer k . The mapping $k \mapsto a(\alpha)$ is affine and the coefficient is $f'(\alpha)$, which is nonzero with overwhelming probability.

The evaluation $b(\alpha)$ is fixed by the relation

$$\bar{h}(z) \cdot \gamma^d - \bar{h}(\gamma z) + \alpha \cdot \frac{a(\alpha)f(\alpha) - f(\alpha)/h(\alpha)}{b(\alpha)} \cdot z^d = N_d(\alpha \cdot z) \cdot f(z^{-1}) \cdot z^d$$

and therefore does not leak any new information. \square